7: Combining Data Sets with `dplyr`

Suppose you have the following two data sets. The first, `df1` has the variables `id_numb` and `xvar`. The second, `df2` has the variables `id` and `yvar`. `id_numb` and `id` serve as identification variables, possibly with duplicates, where observations from the first data set with `id_numb = 1` correspond to observations in the second data set with `id = 1`.

id_numb	xvar
1	16
1	-1
2	11
4	13

id	yvar
1	-1
2	-4
2	0
3	-9

key in df1 key in df2

Mutating Joins

- `left_join()`

dropped left_join(df1, df2), by = c("id_numb" = "id"))

=>

id_numb	xvar	yvar
1	16	-1
1	-1	-1
2	11	-4
2	11	0
4	13	NA

- `right_join()` (id=3 dropped)

dropped right_join(df1, df2)

<=> left_join(df2, df1)

- `inner_join()`

inner_join(df1, df2, by = c("id_numb" = "id"))

=>

id_numb	xvar
1	16
1	-1
2	11
4	13

id	yvar
1	-1
2	-4
2	0
3	-9

- `full_join()`

same as left_join() but adds

id_numb	xvar	yvar
~~4~~ 3	NA	-9

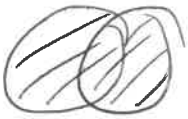

Filtering Joins

- `semi_join()` same syntax

keep anything that has a match in df2

df1, df2 =>

id_numb	xvar
1	16
1	-1
2	11

- `anti_join()` same syntax

keep anything that doesn't have a match in df2.

df1, df2 =>

id_numb	xvar
4	13